

# Mini Project ANN(Shoaib Ghulam Sadar ALi Khan Osama Saleem Huzaifa Abdali)

```
In [2]: ▶ 1 # Part 1 - Data Preprocessing
          2 # Importing the Libraries
          3 import numpy as np
          4 import matplotlib.pyplot as plt
          5 import pandas as pd
```

```
In [3]: ▶ 1 # Importing the dataset
          2 dataset = pd.read_csv('Churn_Modelling.csv')
          3 X = dataset.iloc[:, 3:13].values
          4 y = dataset.iloc[:, 13].values
```

```
In [4]: ▶ 1 # Encoding categorical data
          2 # Encode before splitting because matrix X and independent variable Y mu
          3 # Found two categorical data (country, gender)
          4 # create dummy variables, avoid dummy variable trap
          5 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
          6 labelencoder_X_1 = LabelEncoder()
          7 X[:, 1] = labelencoder_X_1.fit_transform(X[:, 1])
          8 labelencoder_X_2 = LabelEncoder()
          9 X[:, 2] = labelencoder_X_2.fit_transform(X[:, 2])
         10 onehotencoder = OneHotEncoder(categorical_features = [1])
         11 X = onehotencoder.fit_transform(X).toarray()
         12 X = X[:, 1:]
```

D:\Anaconda\lib\site-packages\sklearn\preprocessing\\_encoders.py:371: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

warnings.warn(msg, FutureWarning)

D:\Anaconda\lib\site-packages\sklearn\preprocessing\\_encoders.py:392: DeprecationWarning: The 'categorical\_features' keyword is deprecated in version 0.20 and will be removed in 0.22. You can use the ColumnTransformer instead.

"use the ColumnTransformer instead.", DeprecationWarning)

```
In [5]: ▶ 1 # Splitting the dataset into the Training set and Test set
          2 from sklearn.model_selection import train_test_split
          3
          4 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
```

```
In [8]: 1
        2 # Feature Scaling
        3 # Lots of high computation to ease calculation, we don't want one indepe
        4 from sklearn.preprocessing import StandardScaler
        5 sc = StandardScaler()
        6 X_train = sc.fit_transform(X_train)
        7 X_test = sc.transform(X_test)
```

```
In [9]: 1 # Part 2 - Making the ANN
        2
        3 # Importing the Keras libraries and package
        4 # Sequential module - initialize neural network
        5 # Dense - Layers of ANN
        6 import keras
        7 from keras.models import Sequential
        8 from keras.layers import Dense
        9 from keras.layers import Dropout
```

```
In [10]: 1 # Initialising the ANN
         2 classifier = Sequential()
```

```
In [11]: 1 # Adding the input layer and the first hidden layer with dropout
         2 # Take average of input + output for units/output_dim param in Dense
         3 # input_dim is necessary for the first layer as it was just initialized
         4 classifier.add(Dense(6, input_dim = 11, kernel_initializer = 'glorot_uni
         5 classifier.add(Dropout(p = 0.1))
```

D:\Anaconda\lib\site-packages\ipykernel\_launcher.py:5: UserWarning: Update your `Dropout` call to the Keras 2 API: `Dropout(rate=0.1)`  
"""

```
In [12]: 1 # Adding the second hidden layer with dropout
         2 # doesn't need the input_dim params
         3 # kernel_initializer updates weights
         4 # activation function - rectifier
         5 classifier.add(Dense(6, kernel_initializer = 'glorot_uniform', activatio
         6 classifier.add(Dropout(p = 0.1))
```

D:\Anaconda\lib\site-packages\ipykernel\_launcher.py:6: UserWarning: Update your `Dropout` call to the Keras 2 API: `Dropout(rate=0.1)`

```
In [13]: 1 # Adding the output layer
         2 # dependent variable with more than two categories (3), output_dim needs
         3 classifier.add(Dense(1, kernel_initializer = 'glorot_uniform', activatio
```

```
In [14]: ▶ 1 # Compiling the ANN - applying Stochastic Gradient Descent to whole ANN
          2 # Several different SGD algorithms
          3 # mathematical details based on the loss function
          4 # binary_crossentropy, categorical_crossentropy
          5 classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', met
```

WARNING:tensorflow:From D:\Anaconda\lib\site-packages\tensorflow\python\ops\nn\_impl.py:180: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where

```
In [15]: 1 # Fitting the ANN to the Training Set
        2 # X_train, y_train, Batch size, Epochs (whole training set)
        3 classifier.fit(X_train, y_train, batch_size = 10, nb_epoch = 100)
```

D:\Anaconda\lib\site-packages\ipykernel\_launcher.py:3: UserWarning: The `nb\_epoch` argument in `fit` has been renamed `epochs`.

This is separate from the ipykernel package so we can avoid doing imports until

WARNING:tensorflow:From D:\Anaconda\lib\site-packages\keras\backend\tensorflow\_backend.py:422: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

```
Epoch 1/100
8000/8000 [=====] - 5s 656us/step - loss: 0.5544 -
accuracy: 0.7444
Epoch 2/100
8000/8000 [=====] - 3s 400us/step - loss: 0.4705 -
accuracy: 0.7974
Epoch 3/100
8000/8000 [=====] - 3s 428us/step - loss: 0.4469 -
accuracy: 0.8008
Epoch 4/100
8000/8000 [=====] - 3s 391us/step - loss: 0.4210 -
accuracy: 0.8236
Epoch 5/100
8000/8000 [=====] - 3s 368us/step - loss: 0.4051 -
accuracy: 0.8367
Epoch 6/100
8000/8000 [=====] - 3s 411us/step - loss: 0.3889 -
accuracy: 0.8426
Epoch 7/100
8000/8000 [=====] - 3s 379us/step - loss: 0.3871 -
accuracy: 0.8395
Epoch 8/100
8000/8000 [=====] - 3s 406us/step - loss: 0.3838 -
accuracy: 0.8403
Epoch 9/100
8000/8000 [=====] - 3s 429us/step - loss: 0.3802 -
accuracy: 0.8425
Epoch 10/100
8000/8000 [=====] - 3s 406us/step - loss: 0.3753 -
accuracy: 0.8466
Epoch 11/100
8000/8000 [=====] - 3s 397us/step - loss: 0.3793 -
accuracy: 0.8426
Epoch 12/100
8000/8000 [=====] - 4s 461us/step - loss: 0.3741 -
accuracy: 0.8471
Epoch 13/100
8000/8000 [=====] - 3s 362us/step - loss: 0.3735 -
accuracy: 0.8455
Epoch 14/100
8000/8000 [=====] - 3s 398us/step - loss: 0.3747 -
accuracy: 0.8444
Epoch 15/100
8000/8000 [=====] - 4s 453us/step - loss: 0.3728 -
```

```
accuracy: 0.8474
Epoch 16/100
8000/8000 [=====] - 3s 391us/step - loss: 0.3728 -
accuracy: 0.8469
Epoch 17/100
8000/8000 [=====] - 3s 417us/step - loss: 0.3630 -
accuracy: 0.8514
Epoch 18/100
8000/8000 [=====] - 3s 341us/step - loss: 0.3745 -
accuracy: 0.8444
Epoch 19/100
8000/8000 [=====] - 3s 411us/step - loss: 0.3659 -
accuracy: 0.8499
Epoch 20/100
8000/8000 [=====] - 4s 491us/step - loss: 0.3668 -
accuracy: 0.8491
Epoch 21/100
8000/8000 [=====] - 3s 366us/step - loss: 0.3702 -
accuracy: 0.8466
Epoch 22/100
8000/8000 [=====] - 3s 369us/step - loss: 0.3670 -
accuracy: 0.8489
Epoch 23/100
8000/8000 [=====] - 3s 408us/step - loss: 0.3620 -
accuracy: 0.8481
Epoch 24/100
8000/8000 [=====] - 3s 334us/step - loss: 0.3714 -
accuracy: 0.8472
Epoch 25/100
8000/8000 [=====] - 3s 325us/step - loss: 0.3689 -
accuracy: 0.8476
Epoch 26/100
8000/8000 [=====] - 3s 342us/step - loss: 0.3664 -
accuracy: 0.8494
Epoch 27/100
8000/8000 [=====] - 3s 334us/step - loss: 0.3670 -
accuracy: 0.8505
Epoch 28/100
8000/8000 [=====] - 3s 337us/step - loss: 0.3643 -
accuracy: 0.8494
Epoch 29/100
8000/8000 [=====] - 3s 323us/step - loss: 0.3634 -
accuracy: 0.8497
Epoch 30/100
8000/8000 [=====] - 3s 324us/step - loss: 0.3659 -
accuracy: 0.8499
Epoch 31/100
8000/8000 [=====] - 3s 332us/step - loss: 0.3661 -
accuracy: 0.8484
Epoch 32/100
8000/8000 [=====] - 3s 389us/step - loss: 0.3631 -
accuracy: 0.8503
Epoch 33/100
8000/8000 [=====] - 4s 474us/step - loss: 0.3693 -
accuracy: 0.8435
Epoch 34/100
8000/8000 [=====] - 4s 554us/step - loss: 0.3618 -
```

```
accuracy: 0.8504
Epoch 35/100
8000/8000 [=====] - 4s 457us/step - loss: 0.3630 -
accuracy: 0.8465
Epoch 36/100
8000/8000 [=====] - 3s 429us/step - loss: 0.3582 -
accuracy: 0.8494
Epoch 37/100
8000/8000 [=====] - 3s 403us/step - loss: 0.3621 -
accuracy: 0.8459
Epoch 38/100
8000/8000 [=====] - 3s 402us/step - loss: 0.3649 -
accuracy: 0.8460
Epoch 39/100
8000/8000 [=====] - 4s 450us/step - loss: 0.3687 -
accuracy: 0.8456
Epoch 40/100
8000/8000 [=====] - 4s 454us/step - loss: 0.3644 -
accuracy: 0.8469
Epoch 41/100
8000/8000 [=====] - 4s 442us/step - loss: 0.3587 -
accuracy: 0.8534
Epoch 42/100
8000/8000 [=====] - 4s 458us/step - loss: 0.3621 -
accuracy: 0.8516
Epoch 43/100
8000/8000 [=====] - 3s 423us/step - loss: 0.3582 -
accuracy: 0.8524
Epoch 44/100
8000/8000 [=====] - 4s 442us/step - loss: 0.3641 -
accuracy: 0.8491
Epoch 45/100
8000/8000 [=====] - 3s 386us/step - loss: 0.3622 -
accuracy: 0.8482
Epoch 46/100
8000/8000 [=====] - 4s 464us/step - loss: 0.3642 -
accuracy: 0.8469
Epoch 47/100
8000/8000 [=====] - 3s 374us/step - loss: 0.3568 -
accuracy: 0.8546
Epoch 48/100
8000/8000 [=====] - 3s 410us/step - loss: 0.3615 -
accuracy: 0.8482
Epoch 49/100
8000/8000 [=====] - 3s 407us/step - loss: 0.3626 -
accuracy: 0.8496
Epoch 50/100
8000/8000 [=====] - 3s 344us/step - loss: 0.3595 -
accuracy: 0.8484
Epoch 51/100
8000/8000 [=====] - 3s 386us/step - loss: 0.3618 -
accuracy: 0.8489
Epoch 52/100
8000/8000 [=====] - 3s 390us/step - loss: 0.3607 -
accuracy: 0.8481
Epoch 53/100
8000/8000 [=====] - 3s 387us/step - loss: 0.3568 -
```

```
accuracy: 0.8518
Epoch 54/100
8000/8000 [=====] - 4s 462us/step - loss: 0.3605 -
accuracy: 0.84990s - loss: 0.3600 - accu
Epoch 55/100
8000/8000 [=====] - 4s 465us/step - loss: 0.3640 -
accuracy: 0.8461
Epoch 56/100
8000/8000 [=====] - 3s 403us/step - loss: 0.3595 -
accuracy: 0.8499
Epoch 57/100
8000/8000 [=====] - 3s 347us/step - loss: 0.3598 -
accuracy: 0.8501
Epoch 58/100
8000/8000 [=====] - 4s 459us/step - loss: 0.3591 -
accuracy: 0.8501
Epoch 59/100
8000/8000 [=====] - 4s 456us/step - loss: 0.3596 -
accuracy: 0.8497
Epoch 60/100
8000/8000 [=====] - 3s 404us/step - loss: 0.3598 -
accuracy: 0.8506
Epoch 61/100
8000/8000 [=====] - 4s 485us/step - loss: 0.3580 -
accuracy: 0.8522
Epoch 62/100
8000/8000 [=====] - 4s 560us/step - loss: 0.3576 -
accuracy: 0.8504
Epoch 63/100
8000/8000 [=====] - 4s 526us/step - loss: 0.3612 -
accuracy: 0.8482
Epoch 64/100
8000/8000 [=====] - 4s 443us/step - loss: 0.3600 -
accuracy: 0.8530
Epoch 65/100
8000/8000 [=====] - 4s 457us/step - loss: 0.3562 -
accuracy: 0.8518
Epoch 66/100
8000/8000 [=====] - 4s 475us/step - loss: 0.3614 -
accuracy: 0.8490
Epoch 67/100
8000/8000 [=====] - 4s 484us/step - loss: 0.3598 -
accuracy: 0.8504
Epoch 68/100
8000/8000 [=====] - 4s 444us/step - loss: 0.3573 -
accuracy: 0.8511
Epoch 69/100
8000/8000 [=====] - 4s 479us/step - loss: 0.3582 -
accuracy: 0.8524
Epoch 70/100
8000/8000 [=====] - 4s 450us/step - loss: 0.3623 -
accuracy: 0.8479
Epoch 71/100
8000/8000 [=====] - 4s 453us/step - loss: 0.3631 -
accuracy: 0.8475
Epoch 72/100
8000/8000 [=====] - 4s 476us/step - loss: 0.3624 -
```

```
accuracy: 0.8508
Epoch 73/100
8000/8000 [=====] - 4s 462us/step - loss: 0.3588 -
accuracy: 0.8518
Epoch 74/100
8000/8000 [=====] - 3s 407us/step - loss: 0.3638 -
accuracy: 0.8487
Epoch 75/100
8000/8000 [=====] - 3s 425us/step - loss: 0.3602 -
accuracy: 0.8495
Epoch 76/100
8000/8000 [=====] - 3s 434us/step - loss: 0.3610 -
accuracy: 0.8496
Epoch 77/100
8000/8000 [=====] - 3s 384us/step - loss: 0.3595 -
accuracy: 0.8491
Epoch 78/100
8000/8000 [=====] - 3s 424us/step - loss: 0.3611 -
accuracy: 0.8496
Epoch 79/100
8000/8000 [=====] - 3s 422us/step - loss: 0.3566 -
accuracy: 0.8499
Epoch 80/100
8000/8000 [=====] - 3s 345us/step - loss: 0.3573 -
accuracy: 0.8540
Epoch 81/100
8000/8000 [=====] - 3s 431us/step - loss: 0.3594 -
accuracy: 0.8495
Epoch 82/100
8000/8000 [=====] - 3s 421us/step - loss: 0.3586 -
accuracy: 0.8500
Epoch 83/100
8000/8000 [=====] - ETA: 0s - loss: 0.3602 - accur
acy: 0.85 - 3s 390us/step - loss: 0.3607 - accuracy: 0.8520
Epoch 84/100
8000/8000 [=====] - 3s 387us/step - loss: 0.3560 -
accuracy: 0.8506
Epoch 85/100
8000/8000 [=====] - 3s 428us/step - loss: 0.3571 -
accuracy: 0.8512
Epoch 86/100
8000/8000 [=====] - 3s 417us/step - loss: 0.3588 -
accuracy: 0.8509
Epoch 87/100
8000/8000 [=====] - 3s 378us/step - loss: 0.3551 -
accuracy: 0.8540
Epoch 88/100
8000/8000 [=====] - 3s 403us/step - loss: 0.3538 -
accuracy: 0.8539
Epoch 89/100
8000/8000 [=====] - 3s 368us/step - loss: 0.3553 -
accuracy: 0.8508
Epoch 90/100
8000/8000 [=====] - 3s 398us/step - loss: 0.3597 -
accuracy: 0.8514
Epoch 91/100
8000/8000 [=====] - 3s 423us/step - loss: 0.3565 -
```



```

accuracy: 0.8509
Epoch 92/100
8000/8000 [=====] - 4s 441us/step - loss: 0.3585 -
accuracy: 0.8533
Epoch 93/100
8000/8000 [=====] - 3s 409us/step - loss: 0.3583 -
accuracy: 0.8539
Epoch 94/100
8000/8000 [=====] - 3s 426us/step - loss: 0.3565 -
accuracy: 0.8510
Epoch 95/100
8000/8000 [=====] - 3s 398us/step - loss: 0.3566 -
accuracy: 0.8506
Epoch 96/100
8000/8000 [=====] - 3s 423us/step - loss: 0.3605 -
accuracy: 0.8499
Epoch 97/100
8000/8000 [=====] - 3s 391us/step - loss: 0.3556 -
accuracy: 0.8536
Epoch 98/100
8000/8000 [=====] - 3s 433us/step - loss: 0.3578 -
accuracy: 0.8528
Epoch 99/100
8000/8000 [=====] - 3s 417us/step - loss: 0.3555 -
accuracy: 0.8510
Epoch 100/100
8000/8000 [=====] - 3s 358us/step - loss: 0.3605 -
accuracy: 0.8497

```

Out[15]: <keras.callbacks.callbacks.History at 0x19c3b4ef7f0>

```

In [17]: 1 # Part 3 - Making the predictions and evaluating the model
          2
          3 # Predicting the Test set results
          4 # Training set, see if the new data probability is right
          5 y_pred = classifier.predict(X_test)
          6 y_pred = (y_pred > 0.5)

```

```

In [18]: 1 # Predicting a single new observation
          2 new_prediction = classifier.predict(sc.transform(np.array([[0, 0, 600, 1
          3 new_prediction = (new_prediction > 0.5)

```

D:\Anaconda\lib\site-packages\sklearn\utils\validation.py:595: DataConversionWarning: Data with input dtype int32 was converted to float64 by Standard Scaler.

warnings.warn(msg, DataConversionWarning)

```

In [19]: 1 # Making the Confusion Matrix
          2 from sklearn.metrics import confusion_matrix
          3 cm = confusion_matrix(y_test, y_pred)

```

## Apply the ANN Model

```

In [20]: ▶ 1 # Part 4 - Evaluating, Improving and Tuning the ANN
           2
           3 # Evaluating the ANN
           4 # Keras wrapper and Sci-kit Learn for k-Fold Cross Validation
           5 from keras.wrappers.scikit_learn import KerasClassifier
           6 from sklearn.model_selection import cross_val_score
           7 from keras.models import Sequential
           8 from keras.layers import Dense
           9 def build_classifier():
          10     classifier = Sequential()
          11     classifier.add(Dense(6, input_dim = 11, kernel_initializer = 'glorot_uniform'))
          12     classifier.add(Dense(6, kernel_initializer = 'glorot_uniform', activation = 'relu'))
          13     classifier.add(Dense(1, kernel_initializer = 'glorot_uniform', activation = 'sigmoid'))
          14     classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
          15     return classifier

```

## Apply K-Fold Cross Validation

```

In [21]: ▶ 1 # k-Fold cross validator to check if the real relevant accuracy or the standard deviation
           2 # and where we are in bias-variance tradeoffs
           3 classifier = KerasClassifier(build_fn = build_classifier, batch_size = 1)
           4 accuracies = cross_val_score(estimator = classifier, X = X_train, y = y_train, cv = 5)
           5 mean = accuracies.mean()
           6 variable = accuracies.std()

```

```

In [22]: ▶ 1 # overfitting is when it's trained too much on the training set, less performance on the
           2 # test set, high variance in, Dropout Regularization to reduce overfitting
           3
           4 # Tuning the ANN, parameters learned during training (weights), stay fixed
           5 # parameter tuning best value of these hyperparameters, GridSearchCV with cross_val_score
           6 from keras.wrappers.scikit_learn import KerasClassifier
           7 from sklearn.model_selection import GridSearchCV
           8 from keras.models import Sequential
           9 from keras.layers import Dense
          10 def build_classifier(optimizer):
          11     classifier = Sequential()
          12     classifier.add(Dense(6, input_dim = 11, kernel_initializer = 'glorot_uniform'))
          13     classifier.add(Dense(6, kernel_initializer = 'glorot_uniform', activation = 'relu'))
          14     classifier.add(Dense(1, kernel_initializer = 'glorot_uniform', activation = 'sigmoid'))
          15     classifier.compile(optimizer = optimizer, loss = 'binary_crossentropy', metrics = ['accuracy'])
          16     return classifier
          17

```

```
In [23]: 1 classifier = KerasClassifier(build_fn = build_classifier)
2         parameters = {'batch_size': [25, 32],
3                       'nb_epoch': [100, 500],
4                       'optimizer': ['adam', 'rmsprop']}
5         grid_search = GridSearchCV(estimator = classifier,
6                                    param_grid = parameters,
7                                    scoring = 'accuracy',
8                                    cv = 10)
9         grid_search = grid_search.fit(X_train, y_train)
10        best_parameters = grid_search.best_params_
11        best_accuracy = grid_search.best_score_
```

```
Epoch 1/1
7200/7200 [=====] - 2s 255us/step - loss: 0.6099
- accuracy: 0.7079
Epoch 1/1
7200/7200 [=====] - 2s 334us/step - loss: 0.5758
- accuracy: 0.7539
Epoch 1/1
7200/7200 [=====] - 2s 323us/step - loss: 0.5770
- accuracy: 0.7618
Epoch 1/1
7200/7200 [=====] - 2s 300us/step - loss: 0.5400
- accuracy: 0.7839
Epoch 1/1
7200/7200 [=====] - 2s 313us/step - loss: 0.5648
- accuracy: 0.7835
Epoch 1/1
7200/7200 [=====] - 2s 294us/step - loss: 0.5426
- accuracy: 0.7596
Epoch 1/1
7200/7200 [=====] - 2s 330us/step - loss: 0.5586
```

## Describe The Dataset

In [28]:

```

1 import pandas
2
3 data = pd.read_csv('Churn_Modelling.csv')
4 peek = data.head(20)
5 print(peek)

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
\							
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43
5	6	15574012	Chu	645	Spain	Male	44
6	7	15592531	Bartlett	822	France	Male	50
7	8	15656148	Obinna	376	Germany	Female	29
8	9	15792365	He	501	France	Male	44
9	10	15592389	H?	684	France	Male	27
10	11	15767821	Bearce	528	France	Male	31
11	12	15737173	Andrews	497	Spain	Male	24
12	13	15632264	Kay	476	France	Female	34
13	14	15691483	Chin	549	France	Female	25
14	15	15600882	Scott	635	Spain	Female	35
15	16	15643966	Goforth	616	Germany	Male	45
16	17	15737452	Romeo	653	Germany	Male	58
17	18	15788218	Henderson	549	Spain	Female	24
18	19	15661507	Muldrow	587	Spain	Male	45
19	20	15568982	Hao	726	France	Female	24

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
5	8	113755.78	2	1	0	
6	7	0.00	2	1	1	
7	4	115046.74	4	1	0	
8	4	142051.07	2	0	1	
9	2	134603.88	1	1	1	
10	6	102016.72	2	0	0	
11	3	0.00	2	1	0	
12	10	0.00	2	1	0	
13	5	0.00	2	0	0	
14	7	0.00	2	1	1	
15	3	143129.41	2	0	1	
16	1	132602.88	1	1	0	
17	9	0.00	2	1	1	
18	6	0.00	1	0	0	
19	6	0.00	2	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0

4	79084.10	0
5	149756.71	1
6	10062.80	0
7	119346.88	1
8	74940.50	0
9	71725.73	0
10	80181.12	0
11	76390.01	0
12	26260.98	0
13	190857.79	0
14	65951.65	0
15	64327.26	0
16	5097.67	1
17	14406.41	0
18	158684.81	0
19	54724.03	0

## Dimensions of Your Data

```
In [30]: 1 shape = data.shape
        2 print(shape)

(10000, 14)
```

## Data Type For Each Attribute

```
In [31]: 1 types = data.dtypes
        2 print(types)

RowNumber      int64
CustomerId     int64
Surname        object
CreditScore    int64
Geography      object
Gender         object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object
```

## Descriptive Statistics

```
In [32]: 1 pandas.set_option('display.width', 100)
2 pandas.set_option('precision', 3)
3 description = data.describe()
4 print(description)
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
NumOfProducts \						
count	10000.000	1.000e+04	10000.000	10000.000	10000.000	10000.000
mean	5000.500	1.569e+07	650.529	38.922	5.013	76485.889
std	2886.896	7.194e+04	96.653	10.488	2.892	62397.405
min	1.000	1.557e+07	350.000	18.000	0.000	0.000
25%	2500.750	1.563e+07	584.000	32.000	3.000	0.000
50%	5000.500	1.569e+07	652.000	37.000	5.000	97198.540
75%	7500.250	1.575e+07	718.000	44.000	7.000	127644.240
max	10000.000	1.582e+07	850.000	92.000	10.000	250898.090

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000	10000.000	10000.000	10000.000
mean	0.706	0.515	100090.240	0.204
std	0.456	0.500	57510.493	0.403
min	0.000	0.000	11.580	0.000
25%	0.000	0.000	51002.110	0.000
50%	1.000	1.000	100193.915	0.000
75%	1.000	1.000	149388.247	0.000
max	1.000	1.000	199992.480	1.000

## Correlation Between Attributes

```
In [33]: 1 pandas.set_option('display.width', 100)
2 pandas.set_option('precision', 3)
3 correlations = data.corr(method='pearson')
4 print(correlations)
```

	RowNumber	CustomerId	CreditScore	Age	Tenure
Balance	1.000e+00	0.004	5.840e-03	7.826e-04	-6.495e-03
NumOfProducts	0.007				
RowNumber	4.202e-03	1.000	5.308e-03	9.497e-03	-1.488e-02
CustomerId	0.017				
CreditScore	5.840e-03	0.005	1.000e+00	-3.965e-03	8.419e-04
Age	0.012				
Tenure	7.826e-04	0.009	-3.965e-03	1.000e+00	-9.997e-03
Balance	-0.031				
NumOfProducts	-6.495e-03	-0.015	8.419e-04	-9.997e-03	1.000e+00
HasCrCard	0.013				
IsActiveMember	-9.067e-03	-0.012	6.268e-03	2.831e-02	-1.225e-02
EstimatedSalary	1.000				
Exited	-0.304	0.017	1.224e-02	-3.068e-02	1.344e-02
RowNumber	7.246e-03	0.017	1.224e-02	-3.068e-02	1.344e-02
CustomerId	1.000				
CreditScore	5.987e-04	-0.014	-5.458e-03	-1.172e-02	2.258e-02
Age	0.003				
Tenure	1.204e-02	0.002	2.565e-02	8.547e-02	-2.836e-02
Balance	-0.010				
NumOfProducts	0.010	0.015	-1.384e-03	-7.201e-03	7.784e-03
HasCrCard	-5.988e-03	0.015	-1.384e-03	-7.201e-03	7.784e-03
IsActiveMember	0.014				
EstimatedSalary	-1.657e-02	-0.006	-2.709e-02	2.853e-01	-1.400e-02
Exited	-0.048				

  

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber	5.987e-04	0.012	-0.006	-0.017
CustomerId	-1.403e-02	0.002	0.015	-0.006
CreditScore	-5.458e-03	0.026	-0.001	-0.027
Age	-1.172e-02	0.085	-0.007	0.285
Tenure	2.258e-02	-0.028	0.008	-0.014
Balance	-1.486e-02	-0.010	0.013	0.119
NumOfProducts	3.183e-03	0.010	0.014	-0.048
HasCrCard	1.000e+00	-0.012	-0.010	-0.007
IsActiveMember	-1.187e-02	1.000	-0.011	-0.156
EstimatedSalary	-9.933e-03	-0.011	1.000	0.012
Exited	-7.138e-03	-0.156	0.012	1.000